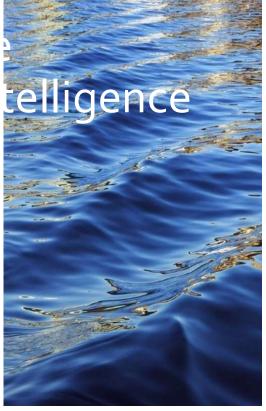


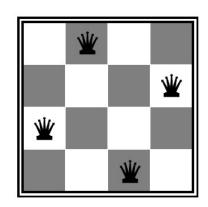
CS 4810 Artificial Intelligence
CS 6810 Topics in Artificial Intelligence





Constraint Satisfaction Problems





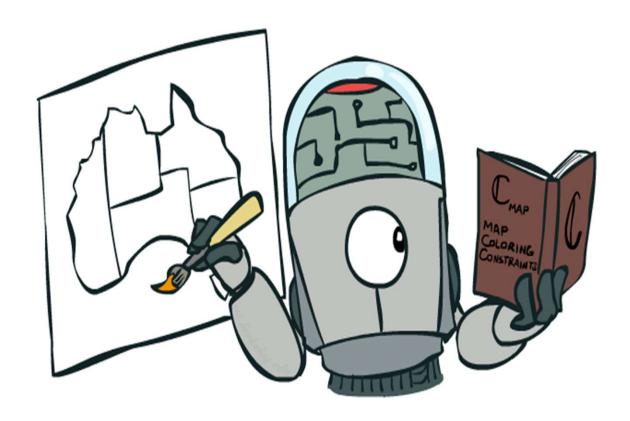
8			4		6			7
						4		
	1					6	5	
5		9		3		7	8	
				7				
	4	8		2		1		3
	5	2					9	5
		1						
3			9		2			5

What is Search For?

- Assumptions about the world: <u>a single agent</u>, <u>deterministic</u> <u>actions</u>, <u>fully observed state</u>, <u>discrete state space</u>
- *Planning*: sequences of actions
 - The path to the goal is the important thing
 - Paths have various costs, depths
 - Heuristics give problem-specific guidance
- *Identification*: assignments to variables
 - The goal itself is important, not the path
 - All paths at the same depth (for some formulat
 - CSPs are specialized for identification problem:

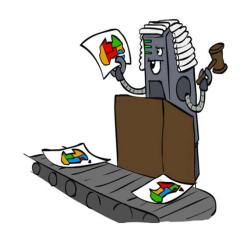


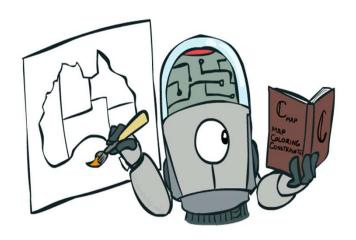
Constraint Satisfaction Problems



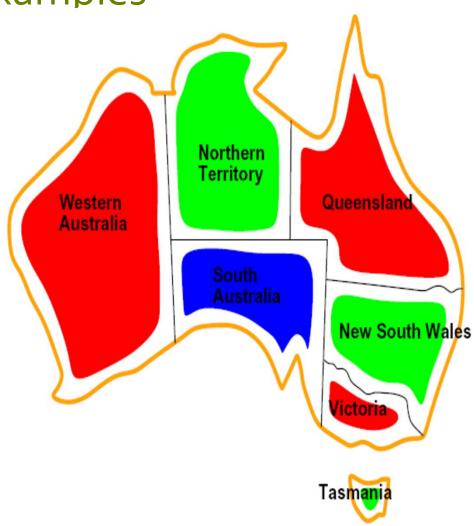
Constraint Satisfaction Problems

- Standard search problems:
 - State is a "black box": arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything
- Constraint satisfaction problems (CSPs):
 - A special subset of search problems
 - State is defined by variables X_i with values from a domain D (sometimes D depends on i)
 - Goal test is a set of constraints specifying allowable combinations of values for subsets of variables
- Simple example of a *formal representation* language
- Allows useful general-purpose algorithms with more power than standard search algorithms





CSP Examples



Example: Map Coloring

• Variables: WA, NT, Q, NSW, V, SA, T

Domains: D = {red, green, blue}

Constraints: adjacent regions must have different colors

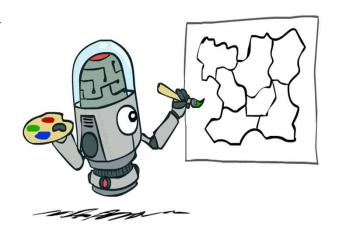
Implicit: $WA \neq NT$

Explicit: $(WA, NT) \in \{(red, green), (red, blue), ...\}$

 Solutions are assignments satisfying all constraints, e.g.:

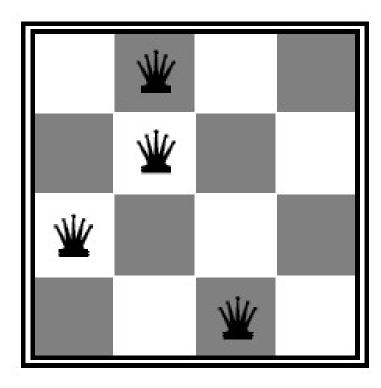
{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green}

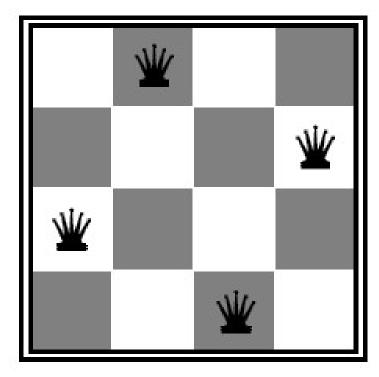




Example: *n*-queens problem

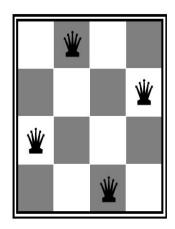
• Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

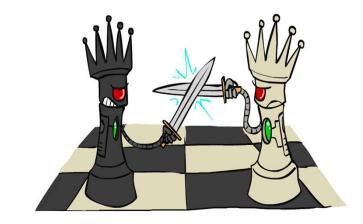




Example: N-Queens

- Formulation 1:
 - Variables: X_{ij} Domains: $\{0,1\}$
 - Constraints





$$\forall i, j, k \ (X_{ij}, X_{ik}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \ (X_{ij}, X_{kj}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \ (X_{ij}, X_{i+k,j+k}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \ (X_{ij}, X_{i+k,j-k}) \in \{(0,0), (0,1), (1,0)\}$$

Example: N-Queens

• Formulation 2:

ullet Variables: Q_k

• Domains: $\{1, 2, 3, ... N\}$

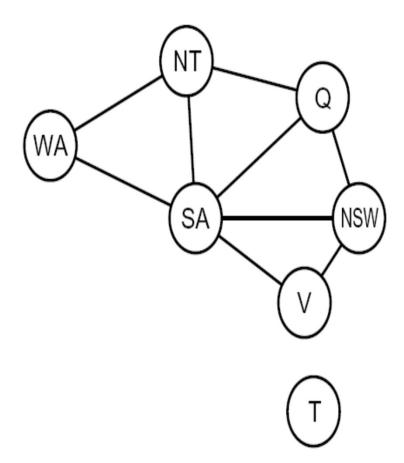
• Constraints:

Implicit: $\forall i, j$ non-threatening (Q_i, Q_j)

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \ldots\}$

• • •

Constraint Graphs

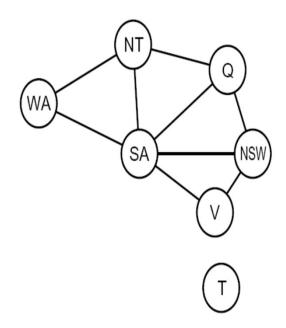


Constraint Graphs

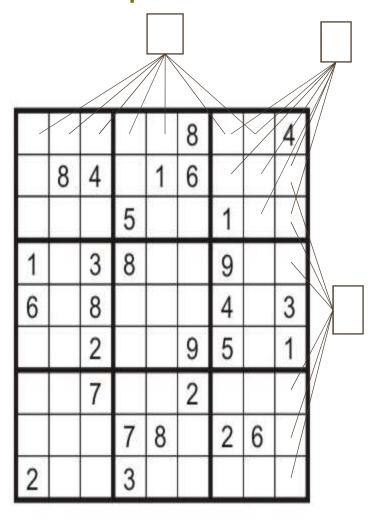
 Binary CSP: each constraint relates (at most) two variables

• Binary constraint graph: nodes are variables, arcs show constraints

 General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!



Example: Sudoku



- Variables:
 - Each (open) square
- Domains:
 - **1**,2,...,9
- Constraints:

9-way alldiff for each column

9-way alldiff for each row

9-way alldiff for each region

(or can have a bunch of pairwise inequality constraints)

Varieties of CSPs

- Discrete Variables
 - Finite domains
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
 - Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Linear constraints solvable, nonlinear undecidable



- E.g., start/end times for Hubble Telescope observations
- Linear constraints solvable in polynomial time by LP methods





Varieties of Constraints

- Varieties of Constraints
 - Unary constraints involve a single variable (equivalent to reducing domains), e.g.:

$$SA \neq green$$

• Binary constraints involve pairs of variables, e.

$$SA \neq WA$$

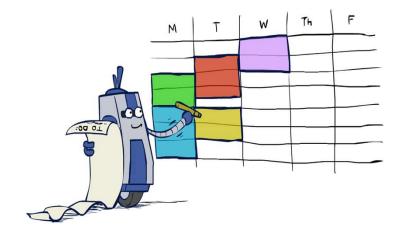
- Higher-order constraints involve 3 or more variables:
- Preferences (soft constraints):

 - E.g., red is better than green
 Often representable by a cost for each variable assignment
 - Gives constrained optimization problems
 - (We'll ignore these until we get to Bayes' nets)



Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- ... lots more!



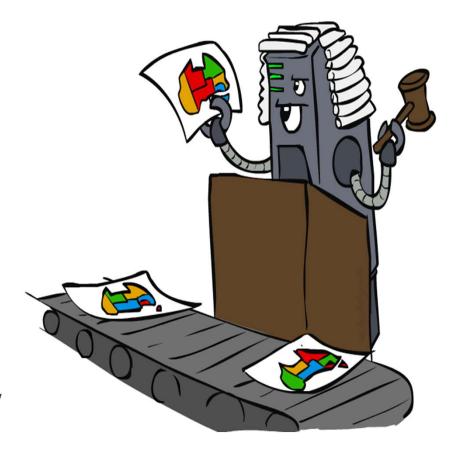
Many real-world problems involve real-valued variables...

Solving CSPs



Standard Search Formulation

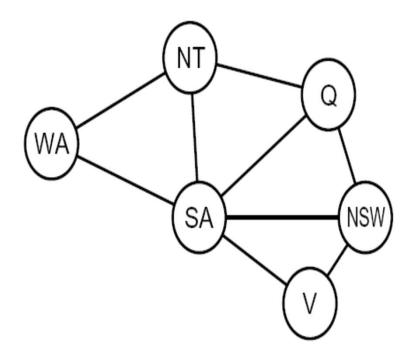
- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment, {}
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with the straightforward, naïve approach, then improve it



Search Methods

• What would BFS do?

What would DFS do?

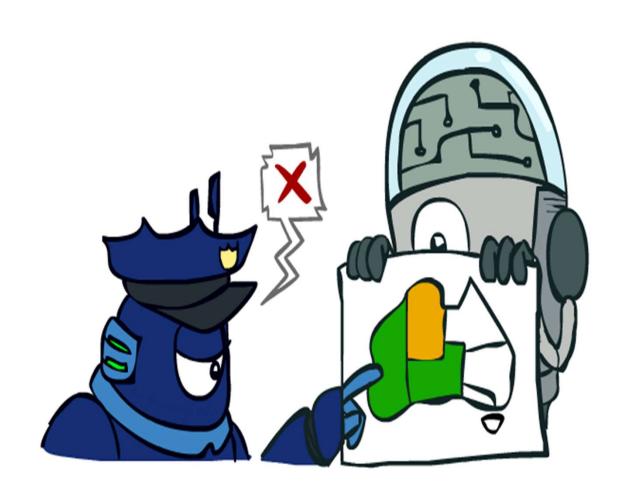


• What problems does naïve search have?

Video of Demo Coloring -- DFS



Backtracking Search



Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to do some computation to check the constraints
 - "Incremental goal test"
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$

