



Constraint Satisfaction Problems




8			4		6			7
						4		
	1					6	5	
5		9		3		7	8	
				7				
	4	8		2		1		3
	5	2					9	
		1						
3			9		2			5

What is Search For?

- Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space
- **Planning**: sequences of actions
 - The path to the goal is the important thing
 - Paths have various costs, depths
 - Heuristics give problem-specific guidance
- **Identification**: assignments to variables
 - The goal itself is important, not the path
 - All paths at the same depth (for some formulations)
 - CSPs are specialized for identification problems





Constraint Satisfaction Problems

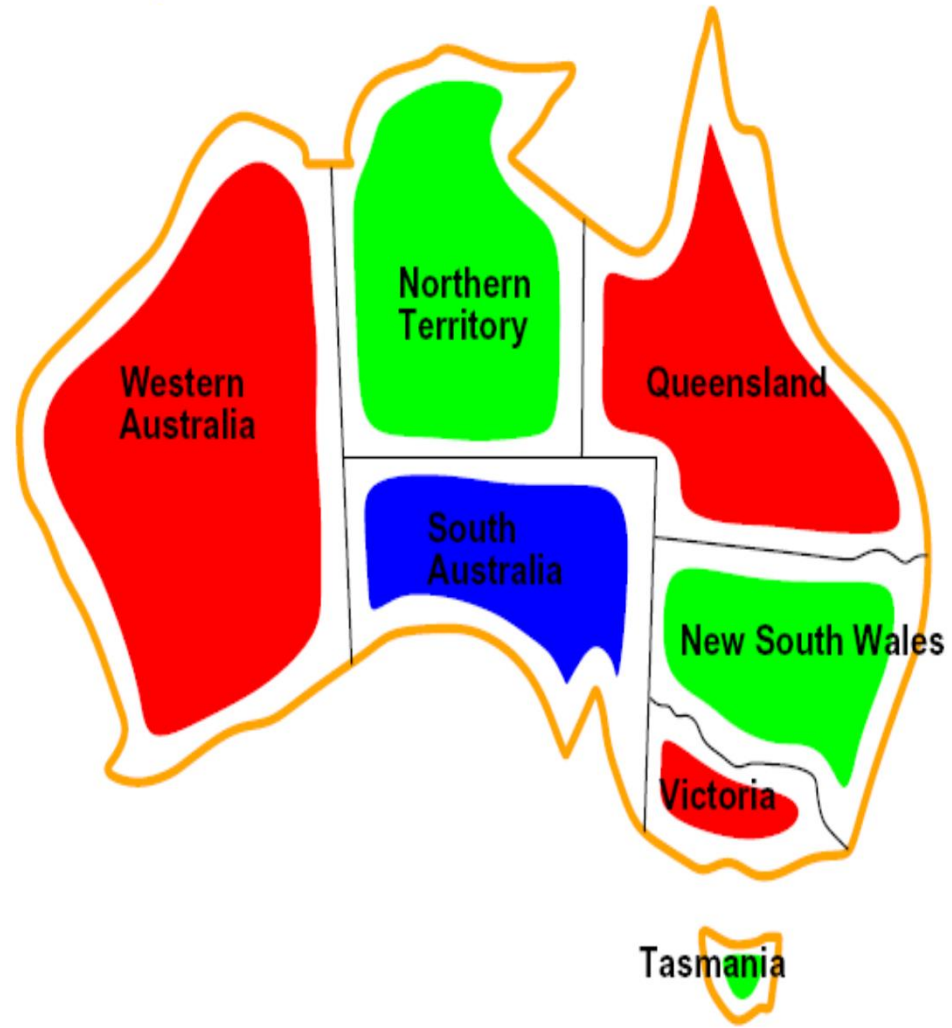


Constraint Satisfaction Problems

- Standard search problems:
 - State is a “black box”: arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything
- Constraint satisfaction problems (CSPs):
 - A special subset of search problems
 - State is defined by **variables X_i** , with values from a **domain D** (sometimes D depends on i)
 - Goal test is a **set of constraints** specifying allowable combinations of values for subsets of variables
- Simple example of a *formal representation language*
- Allows useful general-purpose algorithms with more power than standard search algorithms



CSP Examples



Example: Map Coloring


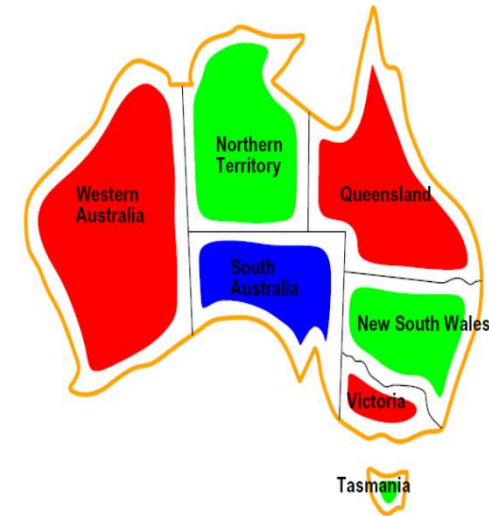
- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D = \{\text{red, green, blue}\}$
- Constraints: adjacent regions must have different colors

Implicit: $WA \neq NT$

Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$



- Solutions are assignments satisfying all constraints, e.g.:

$\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{green}\}$



Example: n -queens problem



- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Example: N-Queens

- Formulation 1:

- Variables: X_{ij}
- Domains: $\{0, 1\}$
- Constraints



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$


$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- Formulation 2:

- Variables: Q_k
- Domains: $\{1, 2, 3, \dots, N\}$
- Constraints:

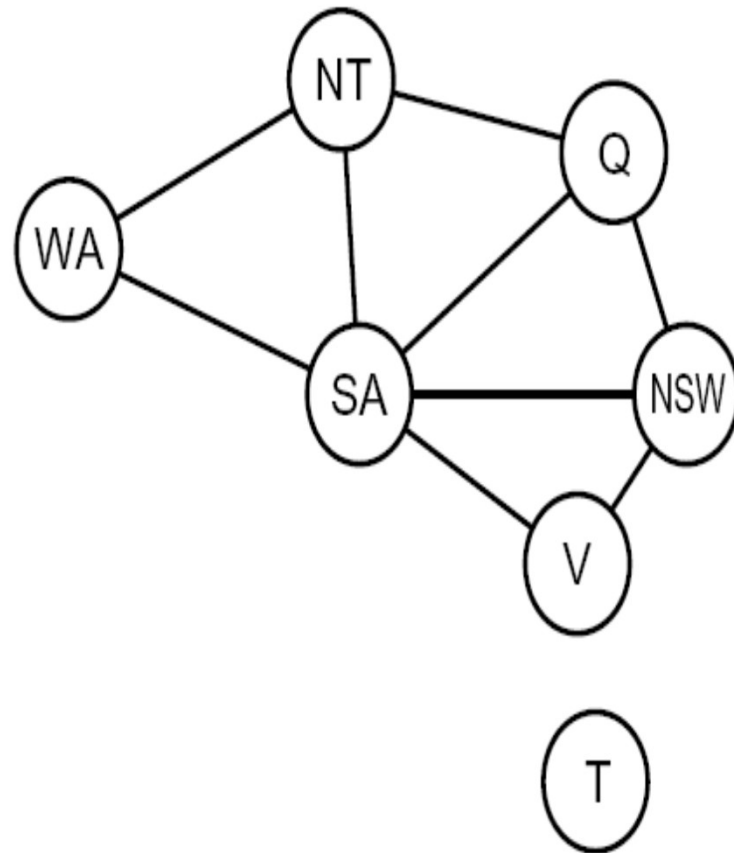


Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$


...

Constraint Graphs




Constraint Graphs

- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!



Example: Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1,2,\dots,9\}$
- Constraints:

9-way alldiff for each column



9-way alldiff for each row

9-way alldiff for each region

(or can have a
bunch of pairwise
inequality
constraints)

Varieties of CSPs

- Discrete Variables
 - Finite domains
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
 - Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Linear constraints solvable, nonlinear undecidable
- Continuous variables
 - E.g., start/end times for Hubble Telescope observations
 - Linear constraints solvable in polynomial time by LP methods



Varieties of Constraints


- Varieties of Constraints
 - Unary constraints involve a single variable (equivalent to reducing domains), e.g.:

$SA \neq \text{green}$

- Binary constraints involve pairs of variables, e.g.


$SA \neq WA$

- Higher-order constraints involve 3 or more variables:
- Preferences (soft constraints):
 - E.g., red is better than green
 - Often representable by a cost for each variable assignment
 - Gives constrained optimization problems
 - (We'll ignore these until we get to Bayes' nets)



Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- ... lots more!




- Many real-world problems involve real-valued variables...

Solving CSPs



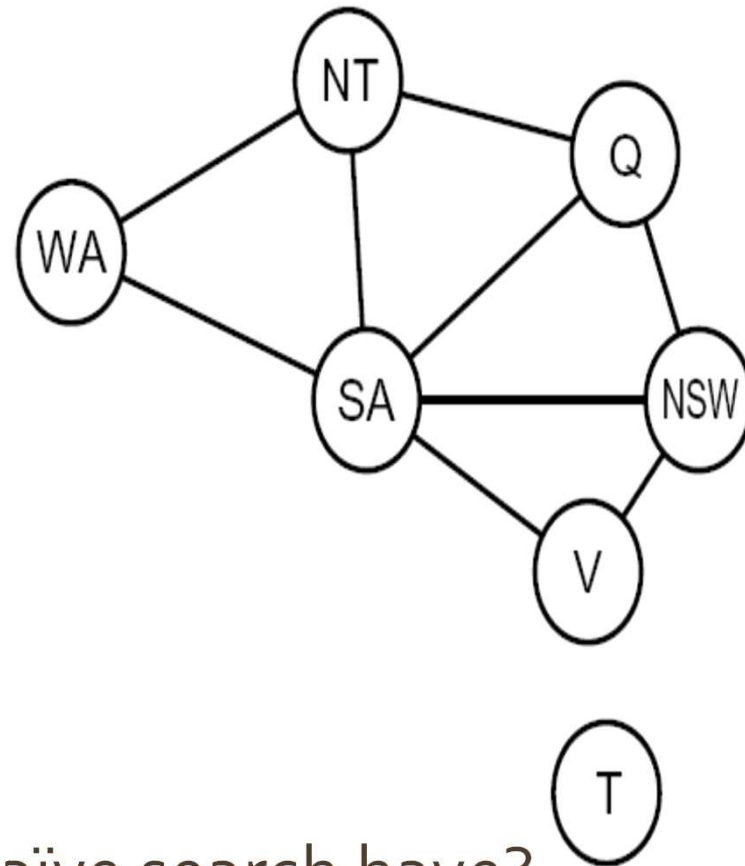
Standard Search Formulation

- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment, $\{\}$
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with the straightforward, naïve approach, then improve it



Search Methods


- What would BFS do?
- What would DFS do?
- What problems does naïve search have?



Video of Demo Coloring -- DFS




Backtracking Search




Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to do some computation to check the constraints
 - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$



Backtracking Example



[Demo: coloring – backtracking Simple]

Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

- Backtracking = DFS + variable-ordering + fail-on-violation
- What are the choice points?


[Demo: coloring – backtracking Complex]

Video of Demo Coloring – Backtracking



Improving Backtracking

- General-purpose ideas give huge gains in speed
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Filtering: Can we detect inevitable failure early?
- Structure: Can we exploit the problem structure?



Filtering



Filtering: Forward Checking

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment



[Demo: coloring -- forward checking]

Which variable should be assigned next?

- **Most constrained variable:**
 - Choose the variable with the fewest legal values
 - A.k.a. **minimum remaining values** (MRV) heuristic



Which variable should be assigned next?

- **Most constrained variable:**
 - Choose the variable with the fewest legal values
 - A.k.a. **minimum remaining values** (MRV) heuristic



Which variable should be assigned next?

- **Most constraining variable:**
 - Choose the variable that imposes the most constraints on the remaining variables



Which variable should be assigned next?

- **Most constraining variable:**
 - Choose the variable that imposes the most constraints on the remaining variables




Given a variable, in which order should its values be tried?

- Choose the **least constraining value**:
 - The value that rules out the fewest values in the remaining variables




Given a variable, in which order should its values be tried?

- Choose the **least constraining value**:
 - The value that rules out the fewest values in the remaining variables




Ordering: Minimum Remaining Values

- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal left values in its domain





- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering




Ordering: Least Constraining Value

- Value Ordering: Least Constraining Value
 - Given a choice of variable, choose the *least constraining value*
 - I.e., the one that rules out the fewest values in the remaining variables
 - Note that it may take some computation to determine this! (E.g., rerunning filtering)
- Why least rather than most?
- Combining these ordering ideas makes 1000 queens feasible



Constraint propagation


- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures



- NT and SA cannot both be blue!
- Constraint propagation** repeatedly enforces constraints *locally*


Arc consistency

- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y




Arc consistency

- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y



Arc consistency


- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y
 - When checking $X \rightarrow Y$, throw out any values of X for which there isn't an allowed value of Y



- If X loses a value, all pairs $Z \rightarrow X$ need to be rechecked

Arc consistency


- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y
 - When checking $X \rightarrow Y$, throw out any values of X for which there isn't an allowed value of Y



- If X loses a value, all pairs $Z \rightarrow X$ need to be rechecked

Arc consistency


- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y
 - When checking $X \rightarrow Y$, throw out any values of X for which there isn't an allowed value of Y



- If X loses a value, all pairs $Z \rightarrow X$ need to be rechecked


Arc consistency

- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y
 - When checking $X \rightarrow Y$, throw out any values of X for which there isn't an allowed value of Y



Arc consistency


- Simplest form of propagation makes each pair of variables **consistent**:
 - $X \rightarrow Y$ is consistent iff for **every** value of X there is **some** allowed value of Y
 - When checking $X \rightarrow Y$, throw out any values of X for which there isn't an allowed value of Y



- Arc consistency detects failure earlier than forward checking
- Can be run before or after each assignment


Tree-structured CSPs

- Certain kinds of CSPs can be solved without resorting to backtracking search!
- *Tree-structured CSP*: constraint graph does not have any loops




Algorithm for tree-structured CSPs

- Choose one variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering




Algorithm for tree-structured CSPs

- Choose one variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering
- **Backward removal phase:** check arc consistency starting from the rightmost node and going backwards



Algorithm for tree-structured CSPs

- Choose one variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering
- **Backward removal phase:** check arc consistency starting from the rightmost node and going backwards
- **Forward assignment phase:** select an element from the domain of each variable going left to right. We are guaranteed that there will be a valid assignment because each arc is consistent




Algorithm for tree-structured CSPs

- If n is the number of variables and m is the domain size, what is the running time of this algorithm?
 - $O(nm^2)$: we have to check arc consistency once for every node in the graph (every node has one parent), which involves looking at pairs of domain values
- What about backtracking search for general CSPs?
 - Worst case $O(m^n)$



Local search for CSPs


- Start with “complete” states, i.e., all variables assigned
- Allow states with unsatisfied constraints
- Attempt to **improve** states by reassigning variable values
- Hill-climbing search:
 - In each iteration, randomly select any conflicted variable and choose value that violates the fewest constraints
 - I.e., attempt to greedily minimize total number of violated constraints



h = number of conflicts

Local search for CSPs

- Start with “complete” states, i.e., all variables assigned
- Allow states with unsatisfied constraints
- Attempt to **improve** states by reassigning variable values
- Hill-climbing search:
 - In each iteration, randomly select any conflicted variable and choose value that violates the fewest constraints
 - I.e., attempt to greedily minimize total number of violated constraints
 - Problem: *local minima*




$h = 1$

Local search for CSPs


- Start with “complete” states, i.e., all variables assigned
- Allow states with unsatisfied constraints
- Attempt to **improve** states by reassigning variable values
- Hill-climbing search:
 - In each iteration, randomly select any conflicted variable and choose value that violates the fewest constraints
 - I.e., attempt to greedily minimize total number of violated constraints
 - Problem: *local minima*
- For more on local search, see ch. 4




N-Queen Iterative Improvement




N-Queen Iterative Improvement




N-Queen Iterative Improvement




N-Queen Iterative Improvement




N-Queen Iterative Improvement




N-Queen Iterative Improvement



N-Queen Iterative Improvement



N-Queen Iterative Improvement



N-Queen Iterative Improvement

