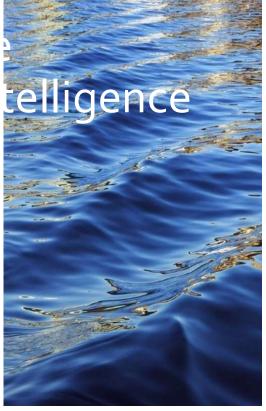
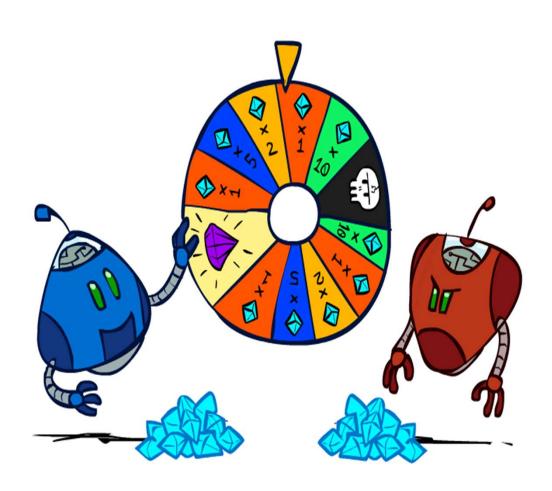


CS 4810 Artificial Intelligence
CS 6810 Topics in Artificial Intelligence

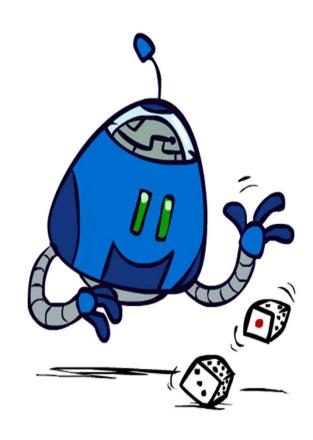




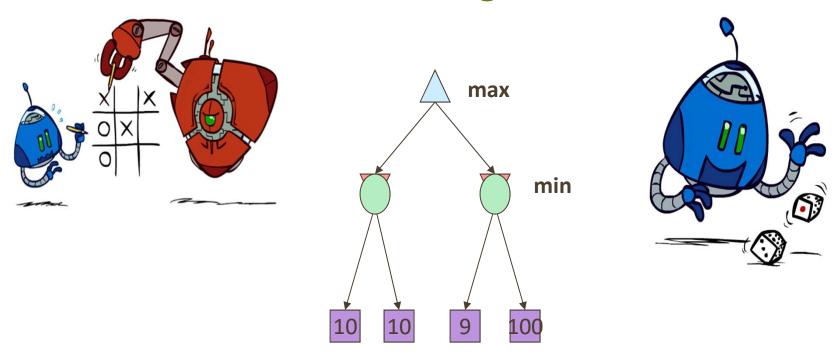
Uncertainty and Utilities



Uncertain Outcomes



Worst-Case vs. Average Case



Idea: Uncertain outcomes controlled by chance, not an adversary!

Expectimax Search

- Why wouldn't we know what the result of an action will be?
 Explicit randomness: rolling dice
 Unpredictable opponents: the ghosts respond

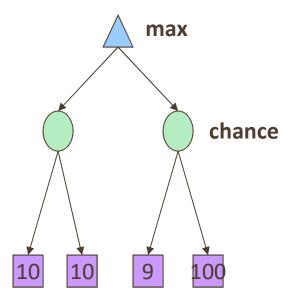
 - randomly
 - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- Expectimax search: compute the average score under optimal play

 Max nodes as in minimax search

 Chance nodes are like min nodes but the outcome

 - is uncertain

 - Calculate their expected utilities
 I.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as Markov Decision Processes



Expectimax Pseudocode

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

def max-value(state):

initialize $v = -\infty$

for each successor of state:

v = max(v, value(successor))

return v

def exp-value(state):

```
initialize v = 0
for each successor of state:
    p = probability(successor)
    v += p * value(successor)
return v
```

Expectimax Pseudocode

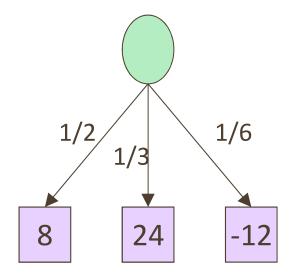
def exp-value(state): initialize v = 0

for each successor of state:

p = probability(successor)

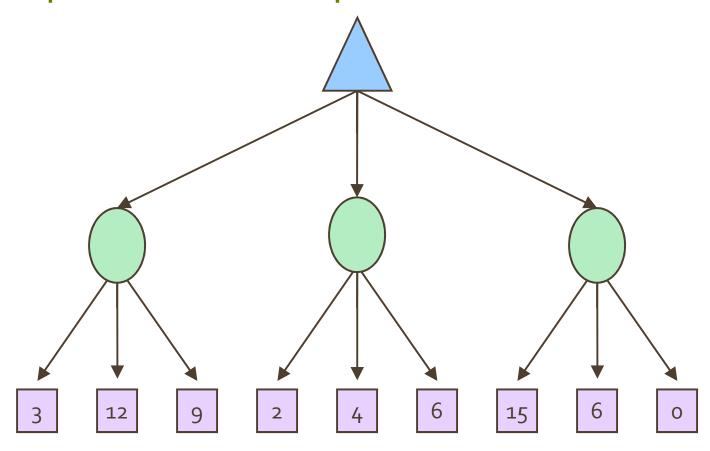
v += p * value(successor)

return v

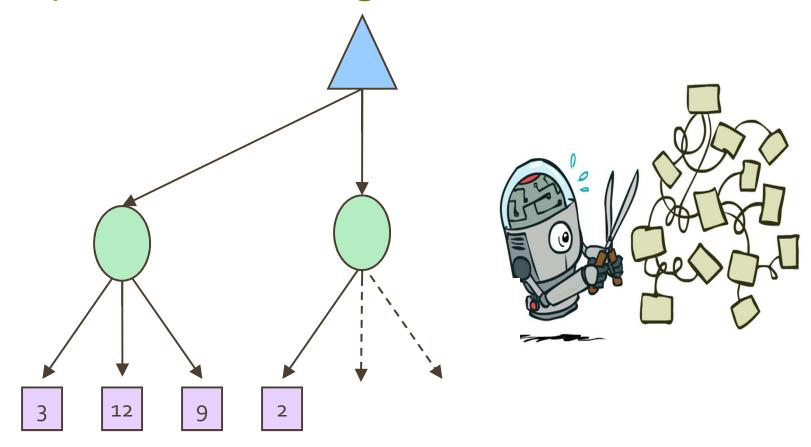


$$v = (1/2)(8) + (1/3)(24) + (1/6)(-12) = 10$$

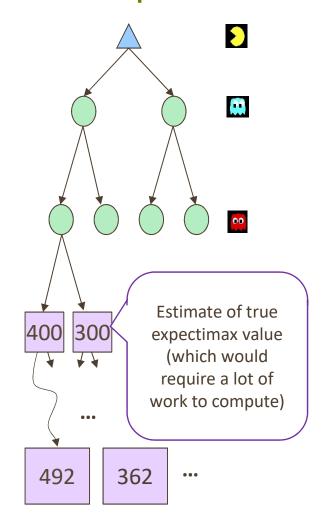
Expectimax Example



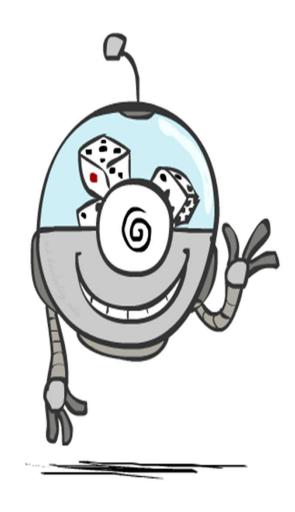
Expectimax Pruning?



Depth-Limited Expectimax

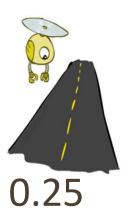


Probabilities



Reminder: Probabilities

- A random variable represents an event whose outcome is unknown
- A probability distribution is an assignment of weights to outcomes



- Example: Traffic on freeway
 - Random variable: T = whether there's traffic
 - Outcomes: T in {none, light, heavy}
 - Distribution: P(T=none) = 0.25, P(T=light) = 0.50, P(T=heavy) = 0.25



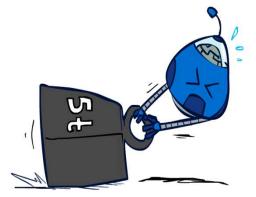
- Some laws of probability (more later):
 - Probabilities are always non-negative
 - Probabilities over all possible outcomes sum to one
- As we get more evidence, probabilities may change:
 - P(T=heavy) = 0.25, P(T=heavy | Hour=8am) = 0.60
 - We'll talk about methods for reasoning and updating probabilities later





Reminder: Expectations

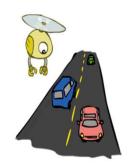
 The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes



Example: How long to get to the airport?

Time: 20 min 30 min 60 min x + x + x 35 min Probability: 0.25 0.50 0.25

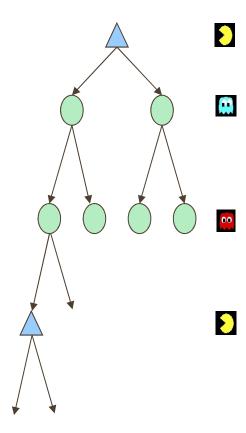






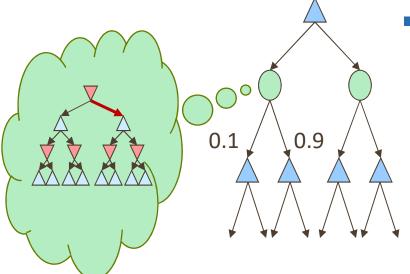
What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
 - Model could be a simple uniform distribution (roll a die)
 - Model could be sophisticated and require a great deal of computation
 - We have a chance node for any outcome out of our control: opponent or environment
 - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



Quiz: Informed Probabilities

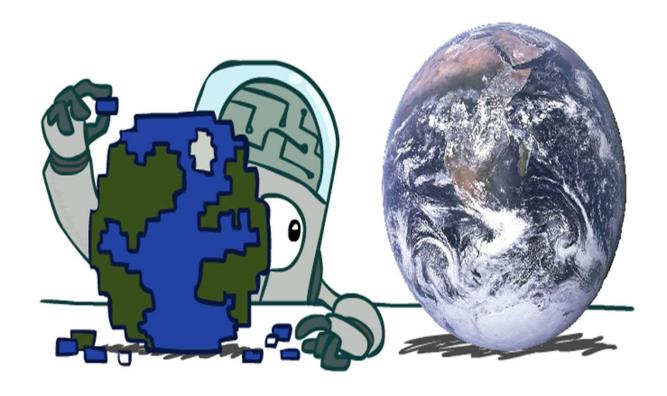
- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



Answer: Expectimax!

- To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
- This kind of thing gets very slow very quickly
- Even worse if you have to simulate your opponent simulating you...
- ... except for minimax, which has the nice property that it all collapses into one game tree

Modeling Assumptions



The Dangers of Optimism and Pessimism

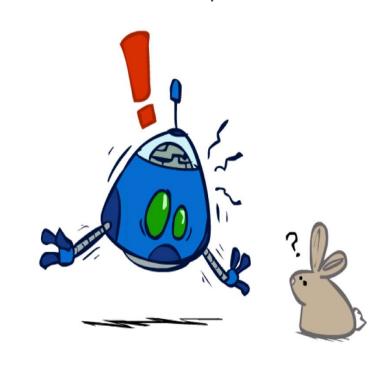
Dangerous Optimism

Assuming chance when the world is adversarial

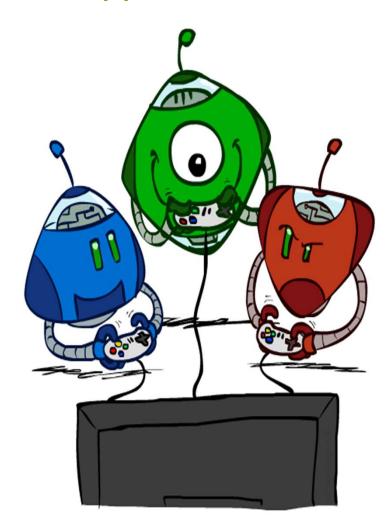


Dangerous Pessimism

Assuming the worst case when it's not likely

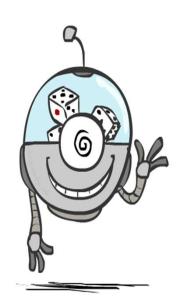


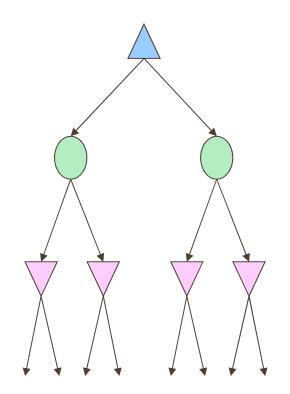
Other Game Types



Mixed Layer Types

- E.g. Backgammon
- Expectiminimax
 - Environment is an extra "random agent" player that moves after each min/max agent
 - Each node computes the appropriate combination of its children











• Dice rolls increase b: 21 possible rolls with 2

- dice
 - Backgammon ≈ 20 legal moves
 - Depth $2 = 20 \times (21 \times 20)^3 = 1.2 \times 10^9$
- As depth increases, probability of reaching a given search node shrinks
 - So usefulness of search is diminished
 - So limiting depth is less damaging
 - But pruning is trickier...
- Historic AI: TDGammon uses depth-2 search + very good evaluation function + reinforcement learning: world-champion level play
- 1st Al world champion in any game!



Multi-Agent Utilities

• What if the game is not zero-sum, or has multiple players?

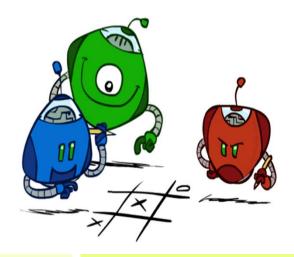


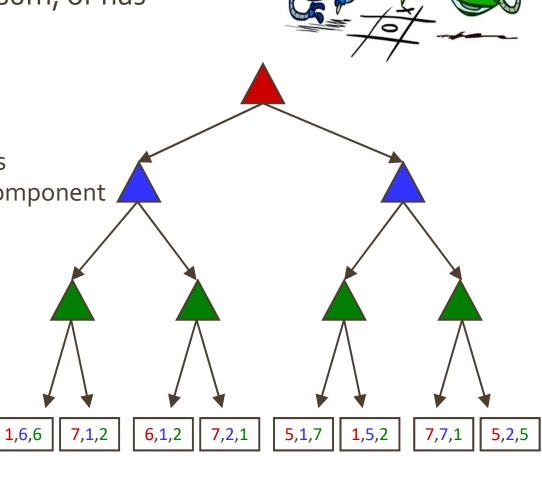
• Terminals have utility tuples

• Node values are also utility tuples

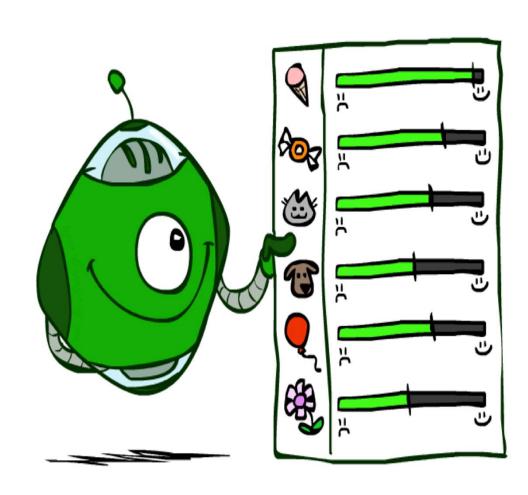
• Each player maximizes its own component

• Can give rise to cooperation and competition dynamically...





Utilities



Maximum Expected Utility

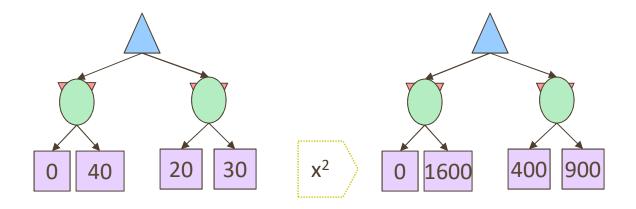
 Why should we average utilities? Why not minimax?

- Principle of maximum expected utility:
 - A rational agent should chose the action that maximizes its expected utility, given its knowledge
- Questions (In next slides):
 - Where do utilities come from?
 - How do we know such utilities even exist?
 - How do we know that averaging even makes sense?
 - What if our behavior (preferences) can't be described by utilities?





What Utilities to Use?



- For worst-case minimax reasoning, terminal function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - We call this insensitivity to monotonic transformations
- For average-case expectimax reasoning, we need *magnitudes* to be meaningful

Utilities

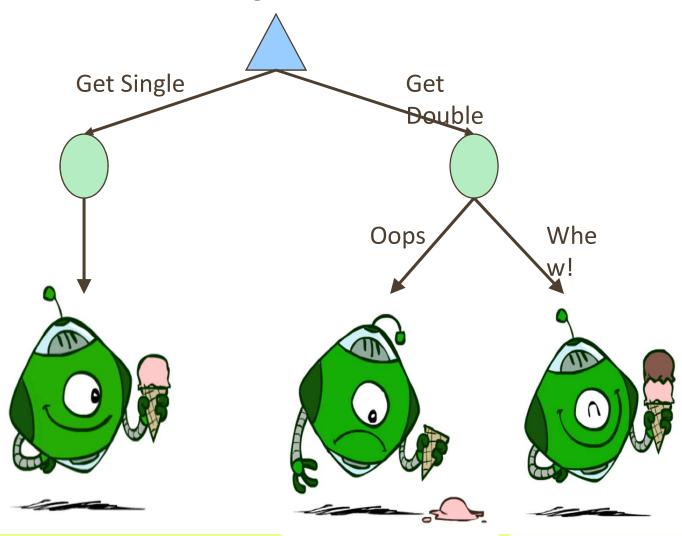
- Utilities are functions from outcomes (states of the world) to real numbers that describe an agent's preferences
- Where do utilities come from?
 - In a game, may be simple (+1/-1)
 - Utilities summarize the agent's goals
 - Theorem: any "rational" preferences can be summarized as a utility function
- We hard-wire utilities and let behaviors emerge
 - Why don't we let agents pick utilities?
 - Why don't we prescribe behaviors?







Utilities: Uncertain Outcomes Getting ice cream



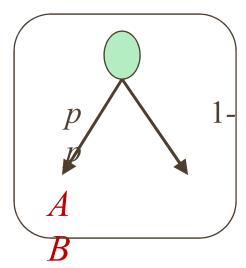
Preferences

- An agent must have preferences among:
 - Prizes: A, B, etc.
 - Lotteries: situations with uncertain prize:L = [p, A; (1-p), B]

A Prize



A Lottery

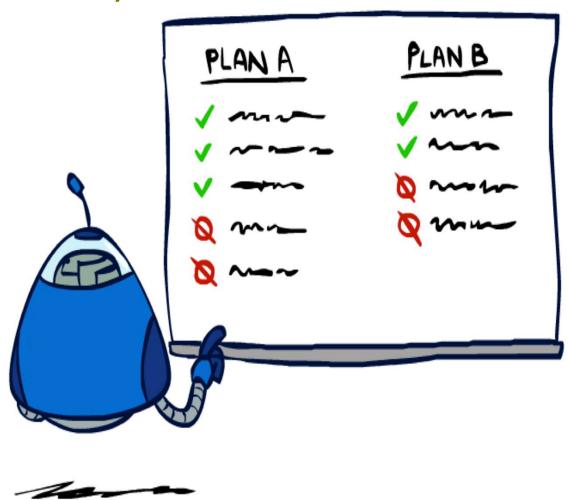


- Notation:
 - Preference: $A \succ B$
 - Indifference: $A \sim B$





Rationality



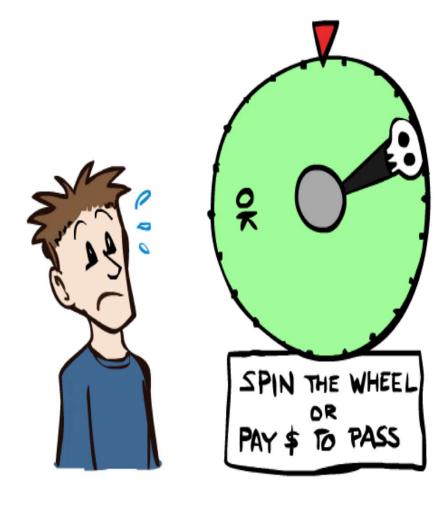
Rational Preferences

 We want some constraints on preferences before we call them rational, such as:

Axiom of Transitivity:
$$(A \succ B) \land (B \succ C) \Longrightarrow (A \succ C)$$

- For example: an agent with intransitive preferences can be induced to give away all of its money
 - If B > C, then an agent with C would pay (say) 1 cent to get B
 - If A > B, then an agent with B would pay (say) 1 cent to get A
 - If C > A, then an agent with A would pay (say) 1 cent to get C

Human Utilities



Utility Scales

- Normalized utilities: $U_{+} = 1.0$, $U_{-} = 0.0$
- Micromorts: one-millionth chance of death, useful for paying to reduce product risks, etc.
- OALYs: quality-adjusted life years, useful for medical decisions involving substantial risk
- Note: behavior is invariant under positive linear transformation

$$U'(x) = k_1 U(x) + k_2$$
 where $k_1 > 0$

• With deterministic prizes only (no lottery choices), only ordinal utility can be determined, i.e., total order on prizes



Human Utilities

- Utilities map states to real numbers. Which number
- Standard approach to assessment (elicitation) of h
 - Compare a prize A to a standard lottery L_D between
 - "best possible prize" υ₊ with probability p
 - "worst possible catastrophe" υ with probability
 - Adjust lottery probability p until indifference:
 - Resulting p is a utility in [0,1]



